

(3D) Computer Vision Tasks: Object Classification, Object Detection, Segmentation (Instance/Semantic/Panoptic), Depth Estimation, CNN-based Regression, Conditional GANs, Neural Rendering, Structure-from-Motion (SfM), 3D Scene Reconstruction, Mesh Generation, Point Cloud Registration, Simulation

Segmentation / Depth Estimation

- Panoptic Segmentation: Fine-tuned **PanopticFCN** for pedestrian segmentation for ADAS/Autonomous Trains.
- Trained depth estimation models with consistency-aware losses and train/test-time augmentations to stress-test robustness.



Neural Rendering

- Developed Neural Rendering pipelines using COLMAP for camera poses, performing point cloud registration, trained NeRFs/Gaussian Splats, and edited Gaussian Volumes in SuperSplats.
- Implemented a Tiny NeRF for forward rendering of road scenes, managing camera extrinsics.




Robotic Simulation with NVIDIA Isaac Sim

- Generate synthetic data by simulating Robot in 3D Gaussian Volumes in Isaac Sim with MobilityGen and the Synthetic Data Generator, **producing occupancy maps** robot navigation, and automated multi-sensor RGB-D and IR thermal camera data collection via the Sensor Processing Graph (SPG), streamed to ROS2 for real-time visualization.
- Converted Fused 3D point clouds into meshes via Marching Cubes, Poisson Surface Reconstruction, and Nvblox algorithms
- Cleaned Meshes in Blender; and used them as proxies for Gaussian Volume in Isaac Sim to generate depth and infrared data,
- Implemented dynamic runtime 3D mesh generation in Unity3D.


Fine-tuning Gaussian Volumes to Fill Blind Spots

- Fine-tuned 3D Gaussian representations by generating VLM scene captions and camera trajectories for novel view synthesis. Used video diffusion models to inpaint under-observed blind spots, distilling synthetic frames back into a second-stage Gaussian optimization to eliminate volumetric artifacts.




Active Learning

- Performed Active Learning for Panoptic Segmentation of pedestrians.
- Scaled test cases every cycle to maintain consistent train-test data distributions.
- Conducted evaluation on unseen test data in every cycle.




Visualizing Image Embedding space for a sequence of images

- Extracted high-dimensional ResNet embeddings to visualize image sequences via t-SNE.
- Applied clustering (DBSCAN, OPTICS, Agglomerative) to automate Scenario Mining.
- Scenario Mining across urban cities, indoor/outdoor depots, countryside, forests, and open fields, identifying edge cases, such as pedestrians in tunnels and diverse pedestrian activities under different lighting.



Contrastive Learning

- Trained a Siamese ResNet-152 backbone using NT-Xent loss to learn robust image representations.
- Mapped feature space to cluster semantically similar images while distancing dissimilar samples.
- Used learned embeddings for downstream tasks, including diverse data sampling.




Uncertainty Estimation

- Standard Strategies: Margin Sampling, Least Confidence, Entropy Criterion.
- Explored Methods: Query-by-Committee (QBC), Monte-Carlo (MC) Dropout, Test-Time Augmentation (TTA), Ensemble Methods.
- QBC with TTA workflow utilizing KL-Divergence to quantify distances between predictive probability distributions.

Fine-tuned Representation Space

Diversity Query Strategy

- Fine-tuned a high-dimensional Embedding space with **Contrastive Learning**
- High-dimensional features are mapped to 2D t-SNE space.
- Diversity Sampling: Used Farthest First/Gonzalez Algorithm for querying samples.



(conditional) GANs

- Developed GANs and conditional GANs (cGANs) from scratch
- Utilized CelebA dataset to generate images with specific attribute combinations.




CNN-based Regression for Autonomous Driving

- Self-Driving Car:** Trained a 5-layer CNN architecture to predict steering angles and perform road-line detection.
- Control Logic: Used predicted steering angles to determine vehicle velocity.




Grad-CAM Analysis

- Performed Grad-CAM analysis to identify pixels driving steering angle predictions.
- Analyzed model behavior in Out-of-Domain scenarios such as rain, fog, and night sky.
- Evaluated model reactions under domain shift.



Image/3D Point Cloud Processing

- Analyzed objects based on location and shape utilizing heatmaps.
- Calculated heatmaps by taking the logarithm of the detected objects.
- Filtered False Positives from inference results before passing to domain experts for labeling.
- Used pixel-wise distributions and Euclidean distance to extract top-3 matching and non-matching samples per image for Contrastive Loss.
- Verified 3D point cloud accuracy by performing Euclidean distance benchmarks.
- Managed 3D data pipelines by converting COLMAP binary (.bin) outputs to .ply and text formats for cross-tool processing.




False Positives

Data Curation

- Semantic Retrieval: Performed data curation for specific scenario retrieval across massive unlabeled datasets.
- Used Zero-shot Transfer for all data curation and retrieval tasks.
- Vision Transformers: Used BLIP/BLIP-2 for image-text semantic retrieval.

Semi-Automatic Annotation and Data Labeling Policy

- Automate the Annotation pipeline for faster processing and training.
- Integrated SAM Docker within CVAT for native AI-assisted labeling.
- Used YOLO, SAM, MaskDINO, Mask2Former, and PanopticFCN for semi-automatic annotation.



- Defined class taxonomy and labeling guidelines, addressed edge cases for consistency.
- Handled RLE to Polygon format conversion and vice-versa.
- Conducted annotation sessions using CVAT.

Evaluation Metrics

- Average Precision (COCO) and Pedestrian Detection Rate
- Inception score (IS) and Fréchet Inception Distance
- PSNR
- Qualitative Analysis
- In-Domain vs Out-of-Domain image evaluation
- Spatial Analysis: Bounding box height vs Width ratio, Segmentation mask area count.



In-Domain Image Sample

Out-of-Domain Image Sample